



Babel: A framework for developing performant and dependable distributed protocols

Pedro Fouto, Pedro Ákos Costa,
Nuno Preguiça, João Leitão

SRDS 2022





Motivation

- Distributed protocol designs are often **simple** (in pseudocode), however implementing them is **error-prone and time consuming**, with lots of low level aspects, such as:
 - Networking
 - Concurrency
 - Multiplexing
 - Protocol interactions
 - ...



Motivation

- We require a framework that **simplifies the development of distributed systems**:
 - **Abstracting low level aspects**: network, concurrency, timers, etc.
 - Without restricting **generability**, allowing implementing any distributed protocol.
 - **Performant** and allowing the implementation of **production-ready** systems.



Babel: Overview

- Java framework providing an **event driven programming model**.
- **Simplifies implementation** of distributed algorithms by trivially translating their specification (pseudo-code).
- Provided abstractions are sufficiently **generic** to support a wide variety of algorithms.
- Capable of supporting **production-ready performant** implementations.



Babel: Overview

- Java framework providing an e
- **Simplifies implementation** of
- specification (pseudo-code).
- Provided abstractions are suffic
- Capable of supporting **product**

upon init do

Send(JOIN, contactNode, myself);

upon Receive(JOIN, newNode) do

trigger addNodeActiveView(newNode)

foreach $n \in \text{activeView}$ and $n \neq \text{newNode}$ **do**

Send(FORWARDJOIN, n , newNode, ARWL, myself)

upon Receive(FORWARDJOIN, newNode, timeToLive, sender) do

if $\text{timeToLive} == 0 \parallel \# \text{activeView} == 1$ **then**

trigger addNodeActiveView(newNode)

else

if $\text{timeToLive} == \text{PRWL}$ **then**

trigger addNodePassiveView(newNode)

$n \leftarrow n \in \text{activeView}$ and $n \neq \text{sender}$

Send(FORWARDJOIN, n , newNode, $\text{timeToLive}-1$, myself)

upon dropRandomElementFromActiveView do

$n \leftarrow n \in \text{activeView}$

Send(DISCONNECT, n , myself)

$\text{activeView} \leftarrow \text{activeView} \setminus \{n\}$

$\text{passiveView} \leftarrow \text{passiveView} \cup \{n\}$

upon addNodeActiveView(node) do

if $\text{node} \neq \text{myself}$ and $\text{node} \notin \text{activeView}$ **then**

if $\text{isfull}(\text{activeView})$ **then**

trigger dropRandomElementFromActiveView

$\text{activeView} \leftarrow \text{activeView} \cup \text{node}$

upon addNodePassiveView(node) do

if $\text{node} \neq \text{myself}$ and $\text{node} \notin \text{activeView}$ and $\text{node} \notin \text{passiveView}$ **then**

if $\text{isfull}(\text{passiveView})$ **then**

$n \leftarrow n \in \text{passiveView}$

$\text{passiveView} \leftarrow \text{passiveView} \setminus \{n\}$

$\text{passiveView} \leftarrow \text{passiveView} \cup \text{node}$

upon Receive(DISCONNECT, peer) do

if $\text{peer} \in \text{activeView}$ **then**

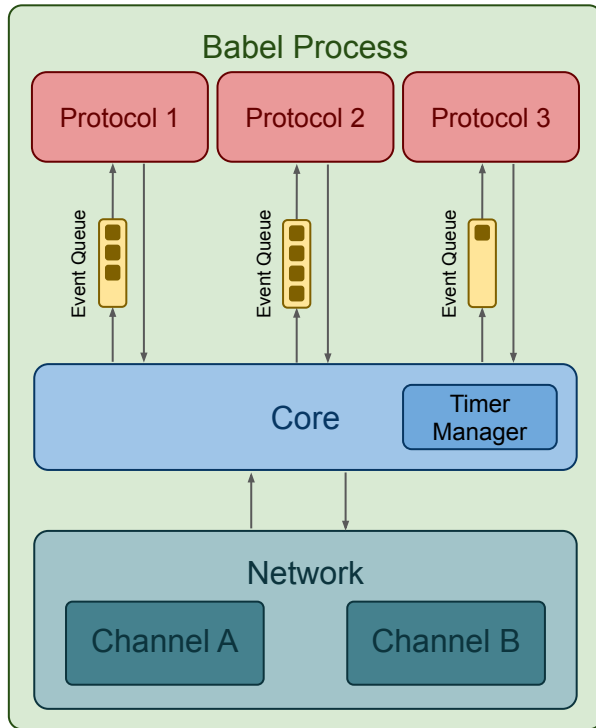
$\text{activeView} \leftarrow \text{activeView} \setminus \{\text{peer}\}$

addNodePassiveView(peer)

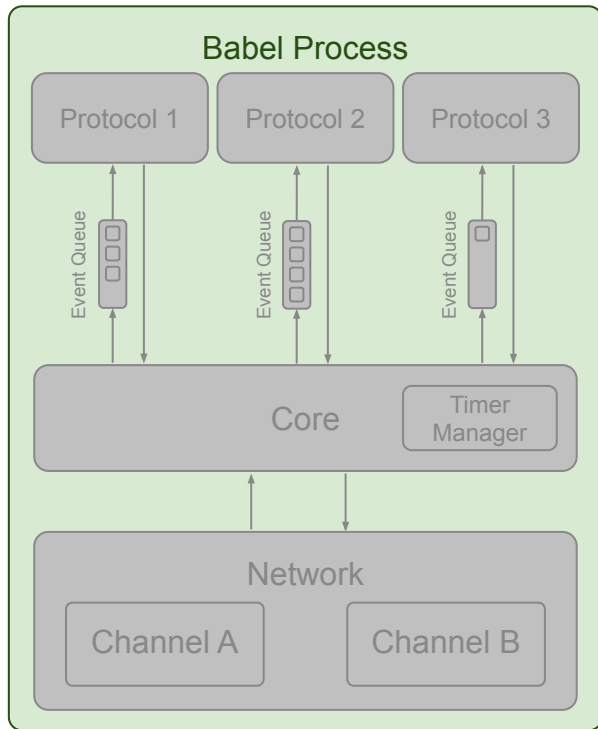
g their

algorithms.

Babel: Architecture

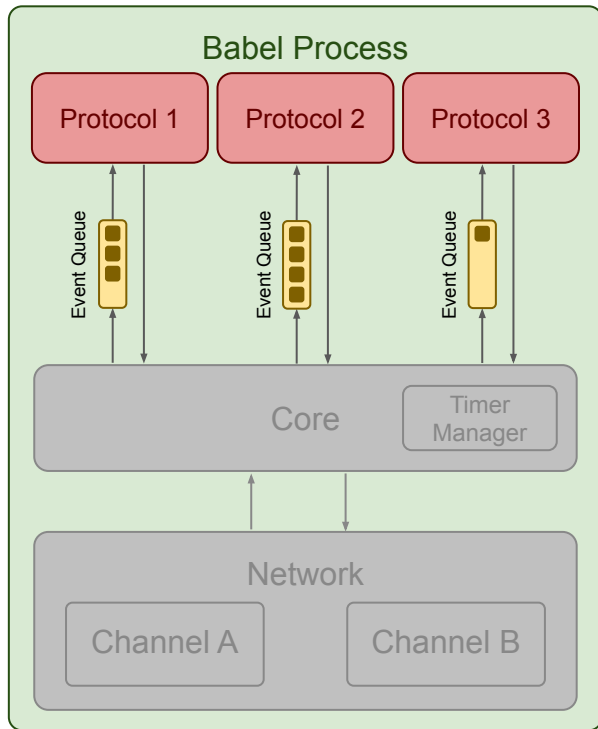


Babel: Architecture




- Each *Babel process* represents a process of the distributed algorithm.

Babel: Architecture



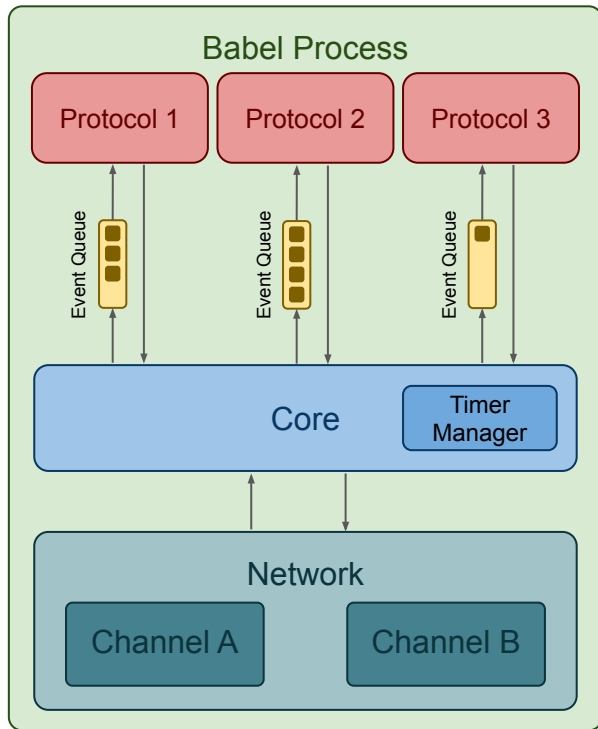
- Each *Babel process* represents a process of the distributed algorithm.
- *Protocols* are implemented by developers and encode the behavior of the algorithm.

Babel: Architecture



- Each *Babel process* represents a process of the distributed algorithm.
- *Protocols* are implemented by developers and encode the behavior of the algorithm.
- The *Core* coordinates and mediates interactions between protocols and processes.

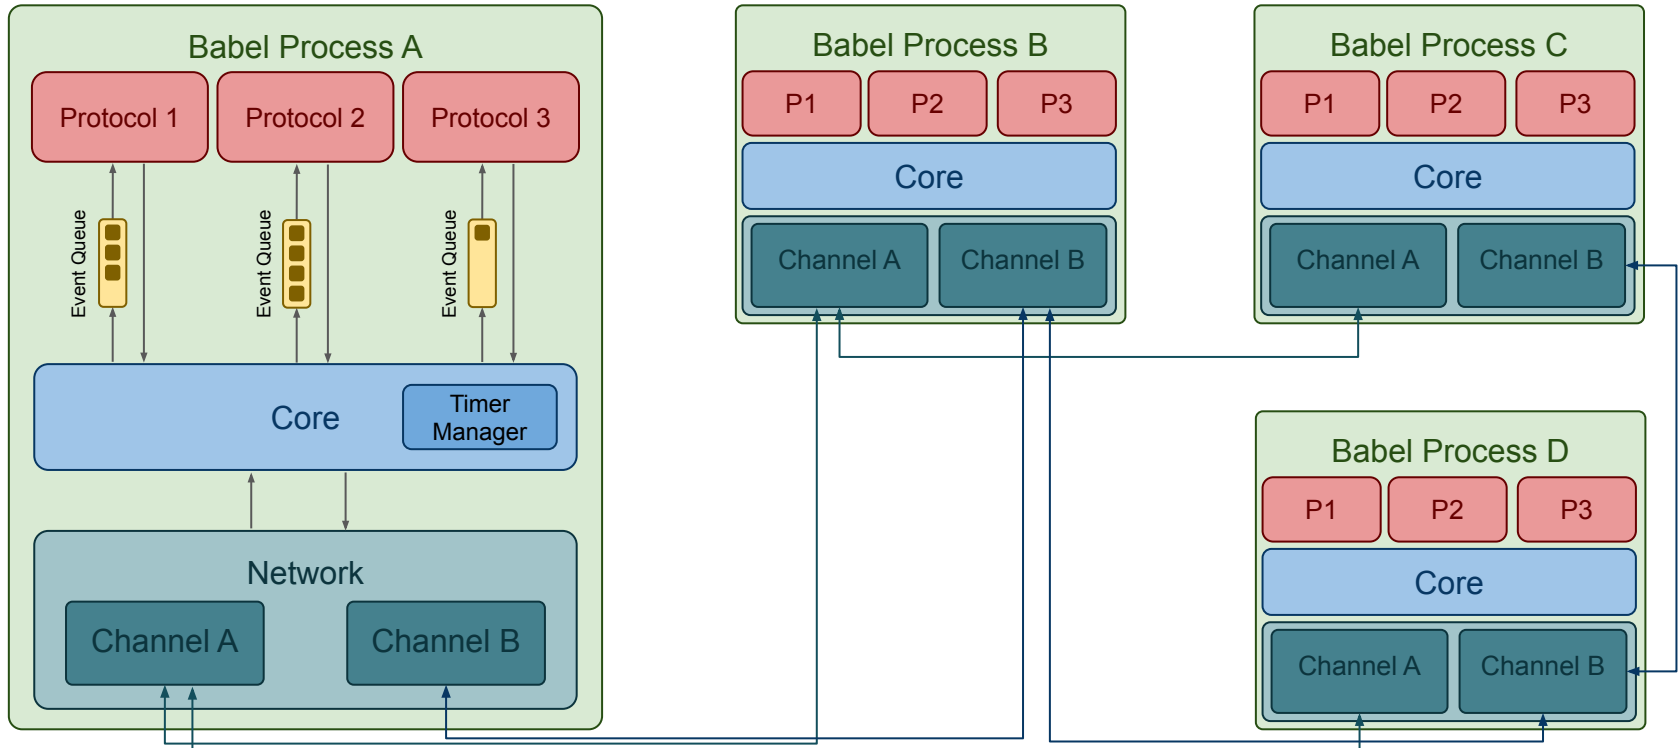
Babel: Architecture



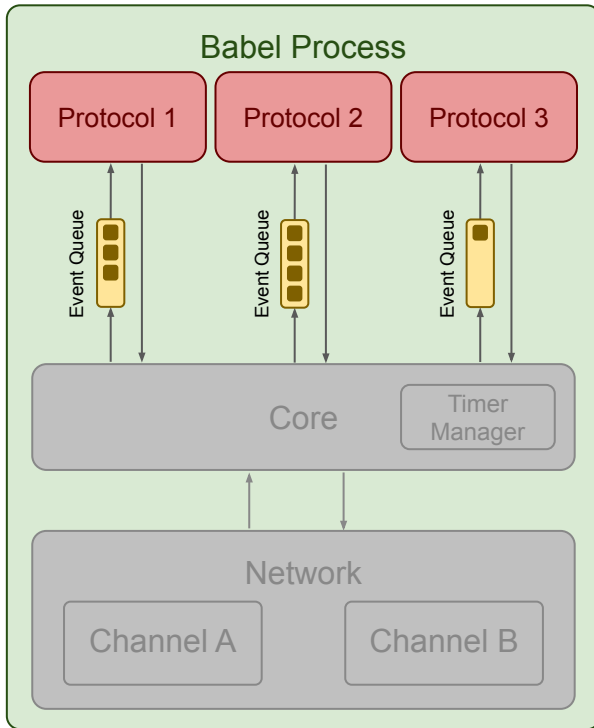
- Each *Babel process* represents a process of the distributed algorithm.
- *Protocols* are implemented by developers and encode the behavior of the algorithm.
- The *Core* coordinates and mediates interactions between protocols and processes.
- *Network channels* abstract communications between processes and provide different behaviours and guarantees.



Babel: Architecture



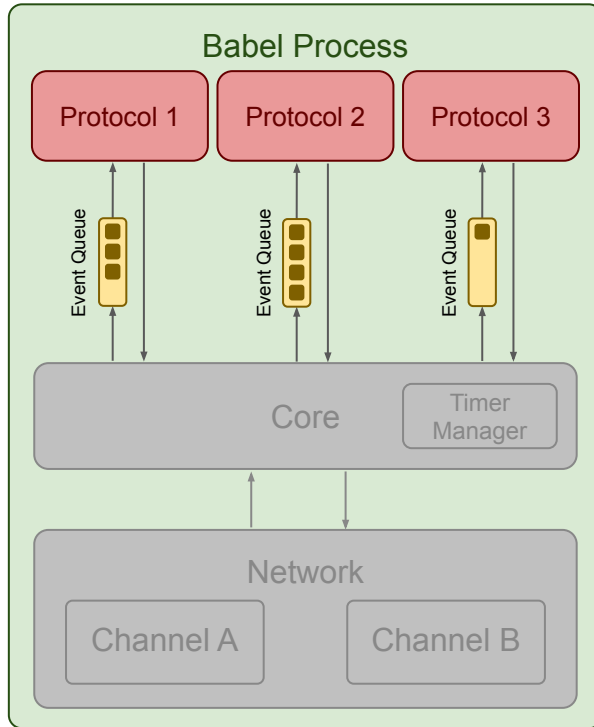
Babel: Architecture



Protocols:

- Encode the behaviour of the distributed algorithm.


Babel: Architecture



Protocols:

- Encode the behaviour of the distributed algorithm.
- Received events are added to the event queue.


Babel: Architecture



Protocols:

- Encode the behaviour of the distributed algorithm.
- Received events are added to the event queue.
- **Dedicated thread** per protocol handles events in **serial** fashion.


Babel: Architecture



Protocols:

- Encode the behaviour of the distributed algorithm.
- Received events are added to the event queue.
- **Dedicated thread** per protocol handles events in **serial** fashion.
- Developer defines and implements **callbacks** for each **event**.


Babel: Architecture



Protocols:

- Encode the behaviour of the distributed algorithm.
- Received events are added to the event queue.
- **Dedicated thread** per protocol handles events in **serial** fashion.
- Developer defines and implements **callbacks** for each **event**.
- **State-machine** that evolves by processing events.


Babel: Architecture



Protocols:

- Encode the behaviour of the distributed algorithm.
- Received events are added to the event queue.
- **Dedicated thread** per protocol handles events in **serial** fashion.
- Developer defines and implements **callbacks** for each **event**.
- **State-machine** that evolves by processing events.
- Communication by **message passing** (no shared state).

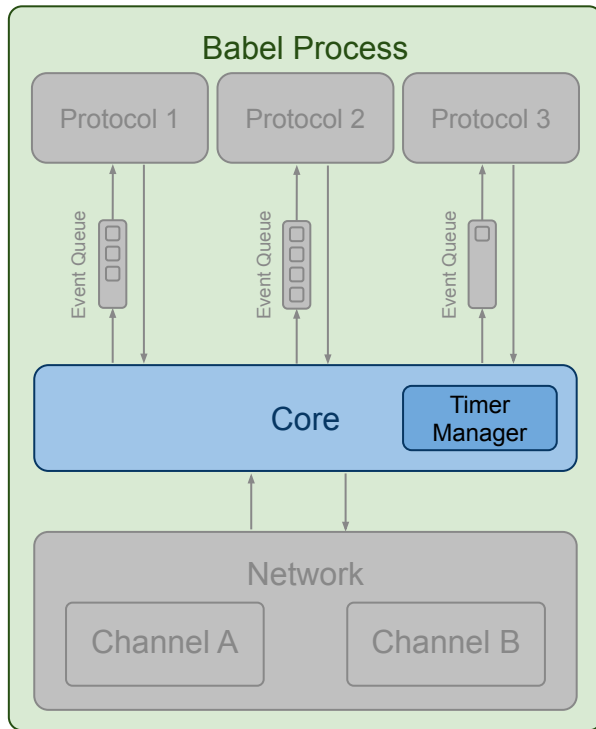
Babel: Architecture



Core:

- Mediates all interactions between protocols.


Babel: Architecture



Core:

- Mediates all interactions between protocols.
 - Exchanges events **between protocols** in the same process.

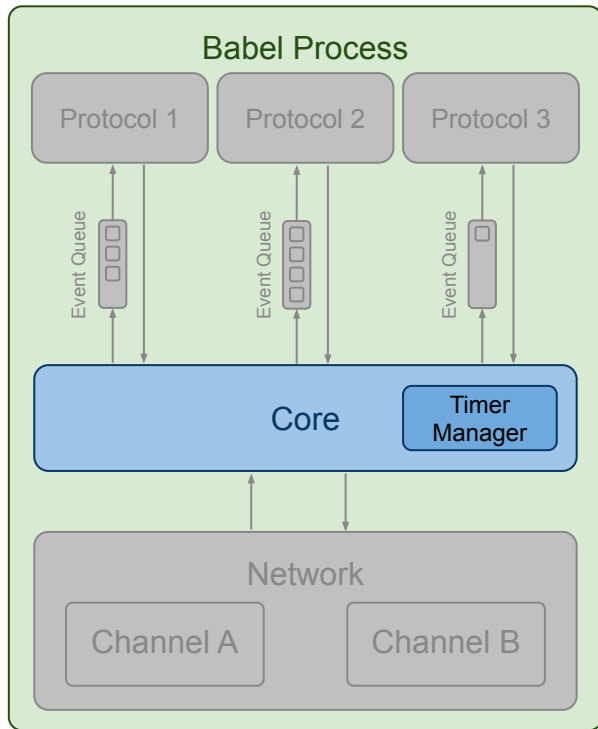
Babel: Architecture



Core:

- Mediates all interactions between protocols.
 - Exchanges events **between protocols** in the same process.
 - Delivers outgoing **network messages** from protocols to the network channel and incoming ones to the correct protocol.

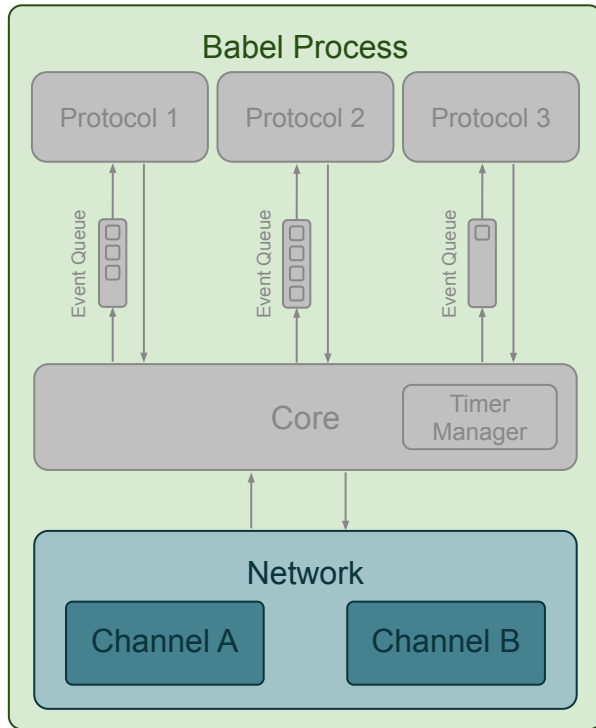
Babel: Architecture



Core:

- Mediates all interactions between protocols.
 - Exchanges events **between protocols** in the same process.
 - Delivers outgoing **network messages** from protocols to the network channel and incoming ones to the correct protocol.
 - Tracks **timers** and delivers timer events to protocols.


Babel: Architecture



Network channels:

- **Abstract complexity** of dealing with networking.

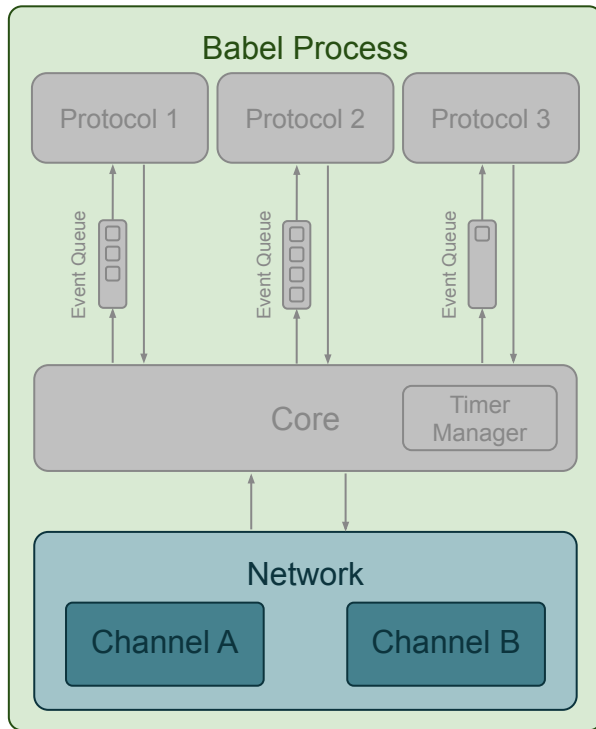
Babel: Architecture



Network channels:

- **Abstract complexity** of dealing with networking.
- Each provides different **guarantees and behaviour**.


Babel: Architecture



Network channels:

- **Abstract complexity** of dealing with networking.
- Each provides different **guarantees and behaviour**.
 - e.g.: explicit and automatic acknowledgement of messages; transparent creation of multiple TCP connections; failure detectors.

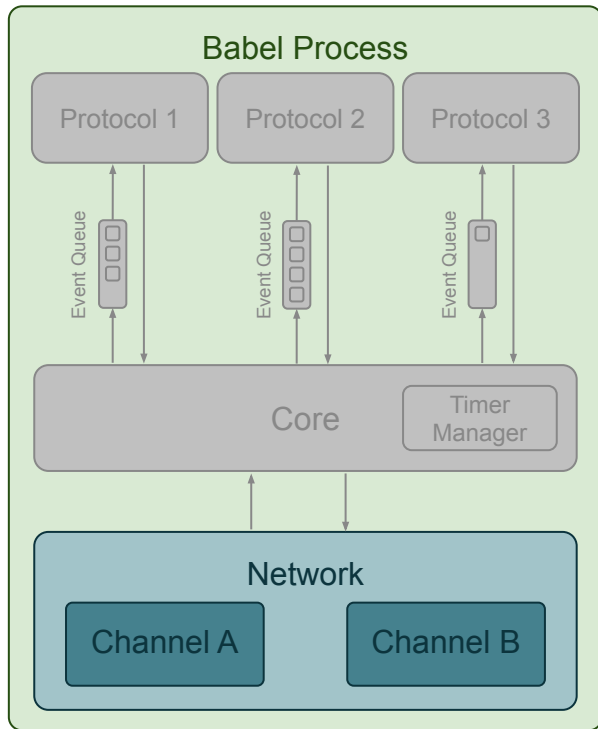
Babel: Architecture



Network channels:

- **Abstract complexity** of dealing with networking.
- Each provides different **guarantees and behaviour**.
 - e.g.: explicit and automatic acknowledgement of messages; transparent creation of multiple TCP connections; failure detectors.
- Developers can **create additional channels** that provides **specific abstractions** for their protocols.

Babel: Architecture



Network channels:

- **Abstract complexity** of dealing with networking.
- Each provides different **guarantees and behaviour**.
 - e.g.: explicit and automatic acknowledgement of messages; transparent creation of multiple TCP connections; failure detectors.
- Developers can **create additional channels** that provides **specific abstractions** for their protocols.
- Protocols can use **multiple channels**, channels can be **shared** by multiple protocols.



Babel: API

```
public class HelloProtocol extends GenericProtocol {  
  
    public HelloProtocol() {  
        super("Hello", 100);  
    }  
  
    public void init (Properties props) {  
        ...  
        sendMessage(new HelloMsg(), contactNode);  
    }  
  
    private void uponHelloMsg(HelloMsg msg, Host from){  
        ...  
    }  
}
```

*Protocols extend an
abstract Java class*



Babel: API

```
public class HelloProtocol extends GenericProtocol {
```

```
    public HelloProtocol() {  
        super("Hello", 100);  
    }
```

```
    public void init (Properties props) {  
        ...  
        sendMessage(new HelloMsg(), contactNode);  
    }
```

```
    private void uponHelloMsg(HelloMsg msg, Host from){  
        ...  
    }  
}
```

*Protocols extend an
abstract Java class*

*Unique identifier for
each protocol*



Babel: API

```
public class HelloProtocol extends GenericProtocol {  
  
    public HelloProtocol() {  
        super("Hello", 100);  
    }  
  
    public void init (Properties props) {  
        ...  
        sendMessage(new HelloMsg(), contactNode);  
    }  
  
    private void uponHelloMsg(HelloMsg msg, Host from){  
        ...  
    }  
}
```

Protocols extend an abstract Java class

Unique identifier for each protocol

Init function initializes the protocol



Babel: API

```
public class HelloProtocol extends GenericProtocol {  
  
    public HelloProtocol() {  
        super("Hello", 100);  
    }  
  
    public void init (Properties props) {  
        ...  
        sendMessage(new HelloMsg(), contactNode);  
    }  
  
    private void uponHelloMsg(HelloMsg msg, Host from) {  
        ...  
    }  
}
```

Protocols extend an abstract Java class

Unique identifier for each protocol

Init function initializes the protocol

Event handlers control protocol flow



Babel: API

```
public class HelloProtocol extends GenericProtocol {  
  
    public HelloProtocol() {  
        super("Hello", 100);  
    }  
  
    public void init (Properties props) {  
        registerMessageHandler(HelloMsg.ID, this::uponHelloMsg);  
        sendMessage(new HelloMsg(), contactNode);  
    }  
  
    private void uponHelloMsg(HelloMsg msg, Host from) {  
        ...  
    }  
}
```

Protocols extend an abstract Java class

Unique identifier for each protocol

Init function initializes the protocol

Event handlers control protocol flow (must be registered)



Babel: API

```
public class HelloProtocol extends GenericProtocol {  
  
    public HelloProtocol() {  
        super("Hello", 100);  
    }  
  
    public void init (Properties props) {  
        registerMessageHandler(HelloMsg.ID, this::uponHelloMsg);  
        sendMessage(new HelloMsg(), contactNode);  
    }  
  
    private void uponHelloMsg(HelloMsg msg, Host from){  
        ...  
    }  
}
```

*Implementing protocols consists
(mostly) in creating and registering
handlers for events*

*Protocols extend an
abstract Java class*

*Unique identifier for
each protocol*

*Init function initializes
the protocol*

*Event handlers control
protocol flow
(must be registered)*



Babel: API

- Protocols handle 3 types of Events:
 - Timers
 - Inter-Protocol Communication
 - Networking



Babel: API

- Protocols handle 3 types of Events:
 - **Timers**
 - Periodic/Timed actions are a common part of distributed applications.
 - Inter-Protocol Communication
 - Networking



Babel: API - Timers

```
public class CounterTimer extends ProtoTimer {  
  
    public static final short ID = 101;  
    int counter;  
  
    public CounterTimer(int initValue) {  
        super(ID);  
        counter = initValue;  
    }  
  
    public void decCounter() {  
        return --counter;  
    }  
}
```

*Extends an abstract
Java class*

Babel: API - Timers

```
public class CounterTimer extends ProtoTimer {  
  
    public static final short ID = 101;  
    int counter;  
  
    public CounterTimer(int initValue) {  
        super(ID);  
        counter = initValue;  
    }  
  
    public void decCounter() {  
        return --counter;  
    }  
}
```

*Extends an abstract
Java class*

*Requires
unique ID*



Babel: API - Timers

```
public class CounterTimer extends ProtoTimer {  
  
    public static final short ID = 101;  
    int counter;  
  
    public CounterTimer(int initValue) {  
        super(ID);  
        counter = initValue;  
    }  
  
    public void decCounter() {  
        return --counter;  
    }  
}
```

*Extends an abstract
Java class*

*Requires
unique ID*

*Can have any
internal logic*



Babel: API - Timers

```
public class CountdownProtocol extends GenericProtocol {

    public CountdownProtocol() {
        super("Countdown", 100);
    }

    @Override
    public void init (Properties props) {
        registerTimerHandler(CounterTimer.ID, this::handleCounter);
        setupPeriodicTimer(new CounterTimer(10), 1000, 300);
    }

    private void handleCounter(CounterTimer timer, long tId){
        doImportantStuff();
        if(timer.decCounter() == 0) cancelTimer(tId);
    }
}
```



Babel: API - Timers

```
public class CountdownProtocol extends GenericProtocol {  
  
    public CountdownProtocol() {  
        super("Countdown", 100);  
    }  
  
    @Override  
    public void init (Properties props) {  
        registerTimerHandler(CounterTimer.ID, this::handleCounter);  
        setupPeriodicTimer(new CounterTimer(10), 1000, 300);  
    }  
  
    private void handleCounter(CounterTimer timer, long tId){  
        doImportantStuff();  
        if(timer.decCounter() == 0) cancelTimer(tId);  
    }  
}
```

*Setup periodic or
"one-shot" timers*



Babel: API - Timers

```
public class CountdownProtocol extends GenericProtocol {

    public CountdownProtocol() {
        super("Countdown", 100);
    }

    @Override
    public void init (Properties props) {
        registerTimerHandler(CounterTimer.ID, this::handleCounter);
        setupPeriodicTimer(new CounterTimer(10), 1000, 300);
    }

    private void handleCounter(CounterTimer timer, long tId){
        doImportantStuff();
        if(timer.decCounter() == 0) cancelTimer(tId);
    }
}
```

*Setup periodic or
"one-shot" timers*

*Setup and register
handlers for timers*



Babel: API - Timers

```
public class CountdownProtocol extends GenericProtocol {

    public CountdownProtocol() {
        super("Countdown", 100);
    }

    @Override
    public void init (Properties props) {
        registerTimerHandler(CounterTimer.ID, this::handleCounter);
        setupPeriodicTimer(new CounterTimer(10), 1000, 300);
    }

    private void handleCounter(CounterTimer timer, long tId){
        doImportantStuff();
        if(timer.decCounter() == 0) cancelTimer(tId);
    }
}
```

*Setup periodic or
"one-shot" timers*

*Setup and register
handlers for timers*



Babel: API - Timers

```
public class CountdownProtocol extends GenericProtocol {  
  
    public CountdownProtocol() {  
        super("Countdown", 100);  
    }  
  
    @Override  
    public void init (Properties props) {  
        registerTimerHandler(CounterTimer.ID, this::handleCounter);  
        setupPeriodicTimer(new CounterTimer(10), 1000, 300);  
    }  
  
    private void handleCounter(CounterTimer timer, long tId){  
        doImportantStuff();  
        if(timer.decCounter() == 0) cancelTimer(tId);  
    }  
}
```

*Setup periodic or
"one-shot" timers*

*Setup and register
handlers for timers*

Cancel timers



Babel: API

- Protocols handle 3 types of Events:
 - Timers
 - Periodic/Timed actions are a common part of distributed applications.
 - Inter-Protocol Communication
 - Networking



Babel: API

- Protocols handle 3 types of Events:
 - Timers
 - Periodic/Timed actions are a common part of distributed applications.
 - **Inter-Protocol Communication**
 - Distributed applications require multiple different protocols to coordinate.
 - One-to-one **requests/replies**
 - One-to-many **notifications**
 - Networking



Babel: API - Inter-protocol communication

```
public class ViewChange extends ProtoNotification {  
  
    public static final short ID = 101;  
    Peer peer;  
  
    public ViewChange(Peer p) {  
        super(ID);  
        peer = p;  
    }  
}
```

*Extends an abstract
Java class*

Babel: API - Inter-protocol communication

```
public class ViewChange extends ProtoNotification {  
  
    public static final short ID = 101;  
    Peer peer;  
  
    public ViewChange(Peer p) {  
        super(ID);  
        peer = p;  
    }  
}
```

*Extends an abstract
Java class*

*Requires
unique ID*



Babel: API - Inter-protocol communication

```
public class ViewChange extends ProtoNotification {  
  
    public static final short ID = 101;  
    Peer peer;  
  
    public ViewChange(Peer p) {  
        super(ID);  
        peer = p;  
    }  
}
```

*Extends an abstract
Java class*

*Requires
unique ID*

*Can have any
internal logic*



Babel: API - Inter-protocol communication

Example:

- **Membership Protocol**
 - Keeps a view with the current membership
- **Dissemination Protocol**
 - Requests current membership and subscribes to changes

Babel: API - Inter-protocol communication

```
public class DisseminationProtocol extends GenericProtocol {  
  
    public DisseminationProtocol() { super("Disseminator", 200); }  
  
    public void init (Properties props) {  
        subscribeNotification(ViewChange.ID, this::onViewChange);  
        registerReplyHandler(ViewReply.ID, this::onViewReply);  
        sendRequest(new ViewRequest(), MembershipProtocol.ID);  
    }  
  
    private void onViewReply(ViewReply reply, short from) {...}  
    private void onViewChange(ViewChange not, short emitter){...}  
}
```



Babel: API - Inter-protocol communication

```
public class DisseminationProtocol extends GenericProtocol {  
  
    public DisseminationProtocol() { super("Disseminator", 200); }  
  
    public void init (Properties props) {  
        subscribeNotification(ViewChange.ID, this::onViewChange);  
        registerReplyHandler(ViewReply.ID, this::onViewReply);  
        sendRequest(new ViewRequest(), MembershipProtocol.ID);  
    }  
  
    private void onViewReply(ViewReply reply, short from) {...}  
    private void onViewChange(ViewChange not, short emitter){...}  
}
```

*Setup and register
handlers*



Babel: API - Inter-protocol communication

```
public class DisseminationProtocol extends GenericProtocol {  
  
    public DisseminationProtocol() { super("Disseminator", 200); }  
  
    public void init (Properties props) {  
        subscribeNotification(ViewChange.ID, this::onViewChange);  
        registerReplyHandler(ViewReply.ID, this::onViewReply);  
        sendRequest(new ViewRequest(), MembershipProtocol.ID);  
    }  
  
    private void onViewReply(ViewReply reply, short from) {...}  
    private void onViewChange(ViewChange not, short emitter){...}  
}
```

*Setup and register
handlers*

*Sending
requests/replies*



Babel: API - Inter-protocol communication

```
public class MembershipProtocol extends GenericProtocol {  
  
    public static final short ID = 300  
  
    public MembershipProtocol() { super("Membership", ID); }  
  
    public void init (Properties props) {  
        registerRequestHandler(ViewRequest.ID, this::onViewRequest);  
    }  
  
    private void onViewRequest(ViewRequest request, short from) {  
        sendReply(new ViewReply(currentView), from);  
    }  
  
    private void onConnectionEstablished(Peer peer){  
        triggerNotification(new ViewChange(peer));  
    }  
}
```



Babel: API - Inter-protocol communication

```
public class MembershipProtocol extends GenericProtocol {  
  
    public static final short ID = 300  
  
    public MembershipProtocol() { super("Membership", ID); }  
  
    public void init (Properties props) {  
        registerRequestHandler(ViewRequest.ID, this::onViewRequest);  
    }  
  
    private void onViewRequest(ViewRequest request, short from) {  
        sendReply(new ViewReply(currentView), from);  
    }  
  
    private void onConnectionEstablished(Peer peer){  
        triggerNotification(new ViewChange(peer));  
    }  
}
```

Replies to received requests



Babel: API - Inter-protocol communication

```
public class MembershipProtocol extends GenericProtocol {  
  
    public static final short ID = 300  
  
    public MembershipProtocol() { super("Membership", ID); }  
  
    public void init (Properties props) {  
        registerRequestHandler(ViewRequest.ID, this::onViewRequest);  
    }  
  
    private void onViewRequest(ViewRequest request, short from) {  
        sendReply(new ViewReply(currentView), from);  
    }  
  
    private void onConnectionEstablished(Peer peer){  
        triggerNotification(new ViewChange(peer));  
    }  
}
```

Replies to received requests

Notifies all subscribers



Babel: API - Inter-protocol communication

```
public class MembershipProtocol extends GenericProtocol {  
  
    public static final short ID = 300  
  
    public MembershipProtocol() { super("Membership", ID); }  
  
    public void init (Properties props) {  
        registerRequestHandler(ViewRequest.ID, this::onViewRequest);  
    }  
  
    private void onViewRequest(ViewRequest request, short from) {  
        sendReply(new ViewReply(currentView), from);  
    }  
  
    private void onConnectionEstablished(Peer peer){  
        triggerNotification(new ViewChange(peer));  
    }  
}
```

Replies to received requests

Notifies all subscribers

No destination protocol required



Babel: API

- Protocols handle 3 types of Events:
 - Timers
 - Periodic/Timed actions are a common part of distributed applications.
 - Inter-Protocol Communication
 - Distributed applications require multiple different protocols to coordinate.
 - One-to-one requests/replies
 - One-to-many notifications
 - Networking



Babel: API

- Protocols handle 3 types of Events:
 - Timers
 - Periodic/Timed actions are a common part of distributed applications.
 - Inter-Protocol Communication
 - Distributed applications require multiple different protocols to coordinate.
 - One-to-one requests/replies
 - One-to-many notifications
 - **Networking**
 - Distributed application need networking to be distributed.



Babel: API - Networking

```
public class PingMsg extends ProtoMessage {
    public final static short MSG_ID = 401;
    private final long timestamp;

    public PingMsg(long timestamp) {
        super(MSG_ID);
        this.timestamp = timestamp;
    }

    public static ISerializer<PingMsg> serializer = new ISerializer<>() {
        public void serialize(PingMsg msg, ByteBuf out) {
            out.writeLong(msg.timestamp);
        }

        public PingMsg deserialize(ByteBuf in){
            return new PingMsg(in.readLong());
        }
    }
}
```

*Network messages
extend generic class*



Babel: API - Networking

```
public class PingMsg extends ProtoMessage {
    public final static short MSG_ID = 401;
    private final long timestamp;

    public PingMsg(long timestamp) {
        super(MSG_ID);
        this.timestamp = timestamp;
    }

    public static ISerializer<PingMsg> serializer = new ISerializer<>() {
        public void serialize(PingMsg msg, ByteBuf out) {
            out.writeLong(msg.timestamp);
        }

        public PingMsg deserialize(ByteBuf in){
            return new PingMsg(in.readLong());
        }
    }
}
```

*Network messages
extend generic class*

*Can have any internal
logic*



Babel: API - Networking

```
public class PingMsg extends ProtoMessage {
    public final static short MSG_ID = 401;
    private final long timestamp;

    public PingMsg(long timestamp) {
        super(MSG_ID);
        this.timestamp = timestamp;
    }

    public static ISerializer<PingMsg> serializer = new ISerializer<>() {
        public void serialize(PingMsg msg, ByteBuf out) {
            out.writeLong(msg.timestamp);
        }

        public PingMsg deserialize(ByteBuf in) {
            return new PingMsg(in.readLong());
        }
    }
}
```

*Network messages
extend generic class*

*Can have any internal
logic*

*Require a serializer to
encode and decode*



Babel: API - Networking

```
public class PingMsg extends ProtoMessage {
    public final static short MSG_ID = 401;
    private final long timestamp;

    public PingMsg(long timestamp) {
        super(MSG_ID);
        this.timestamp = timestamp;
    }

    public static ISerializer<PingMsg> serializer = new ISerializer<>() {
        public void serialize(PingMsg msg, ByteBuf out) {
            out.writeLong(msg.timestamp);
        }

        public PingMsg deserialize(ByteBuf in) {
            return new PingMsg(in.readLong());
        }
    }
}
```

*Network messages
extend generic class*

*Can have any internal
logic*

*Require a serializer to
encode and decode*



Babel: API - Networking

```
public class PingMsg extends ProtoMessage {
    public final static short MSG_ID = 401;
    private final long timestamp;

    public PingMsg(long timestamp) {
        super(MSG_ID);
        this.timestamp = timestamp;
    }

    public static ISerializer<PingMsg> serializer = new ISerializer<>() {
        public void serialize(PingMsg msg, ByteBuf out) {
            out.writeLong(msg.timestamp);
        }

        public PingMsg deserialize(ByteBuf in) {
            return new PingMsg(in.readLong());
        }
    }
}
```

*Network messages
extend generic class*

*Can have any internal
logic*

*Require a serializer to
encode and decode*



Babel: API - Networking

Example:

- **Ping Pong Protocol**
 - Chooses a random peer to send a Ping message
 - Peer responds with Pong message
 - Repeat



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {
    List<Host> peers;

    public PingPongProto() { super("PingPong", 400); }

    public void init (Properties props) {
        int cId = createChannel(TCPChannel.Name, props);
        registerMessageSerializer(cId, PingMsg.ID, PingMsg.serializer);
        registerMessageHandler(cId, PingMsg.ID, this::uponPing);
        registerChannelEventHandler(cId, OutConnUp.EVENT_ID,
            this::uponOutConnUp);
        openConnection(random(peers));
    }
}
```




Babel: API - Networking

Create channel

```
public class PingPongProto extends GenericProtocol {
    List<Host> peers;

    public PingPongProto() { super("PingPong", 400); }

    public void init (Properties props) {
        int cId = createChannel(TCPChannel.Name, props);
        registerMessageSerializer(cId, PingMsg.ID, PingMsg.serializer);
        registerMessageHandler(cId, PingMsg.ID, this::uponPing);
        registerChannelEventHandler(cId, OutConnUp.EVENT_ID,
            this::uponOutConnUp);
        openConnection(random(peers));
    }
}
```



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {
    List<Host> peers;

    public PingPongProto() { super("PingPong", 400); }

    public void init (Properties props) {
        int cId = createChannel(TCPChannel.Name, props);
        registerMessageSerializer(cId, PingMsg.ID, PingMsg.serializer);
        registerMessageHandler(cId, PingMsg.ID, this::uponPing);
        registerChannelEventHandler(cId, OutConnUp.EVENT_ID,
            this::uponOutConnUp);
        openConnection(random(peers));
    }
}
```

Create channel

*Register serializer for
each message*



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {
    List<Host> peers;

    public PingPongProto() { super("PingPong", 400); }

    public void init (Properties props) {
        int cId = createChannel(TCPChannel.Name, props);
        registerMessageSerializer(cId, PingMsg.ID, PingMsg.serializer);
        registerMessageHandler(cId, PingMsg.ID, this::uponPing);
        registerChannelEventHandler(cId, OutConnUp.EVENT_ID,
            this::uponOutConnUp);
        openConnection(random(peers));
    }
}
```

Create channel

*Register serializer for
each message*

*Register handler for
each message*



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {
    List<Host> peers;

    public PingPongProto() { super("PingPong", 400); }

    public void init (Properties props) {
        int cId = createChannel(TCPChannel.Name, props);
        registerMessageSerializer(cId, PingMsg.ID, PingMsg.serializer);
        registerMessageHandler(cId, PingMsg.ID, this::uponPing);
        registerChannelEventHandler(cId, OutConnUp.EVENT_ID,
            this::uponOutConnUp);
        openConnection(random(peers));
    }
}
```

Create channel

*Register serializer for
each message*

*Register handler for
each message*

*Register other
channel events*



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {
    List<Host> peers;

    public PingPongProto() { super("PingPong", 400); }

    public void init (Properties props) {
        int cId = createChannel(TCPChannel.Name, props);
        registerMessageSerializer(cId, PingMsg.ID, PingMsg.serializer);
        registerMessageHandler(cId, PingMsg.ID, this::uponPing);
        registerChannelEventHandler(cId, OutConnUp.EVENT_ID,
            this::uponOutConnUp);
        openConnection(random(peers));
    }
}
```

Create channel

*Register serializer for
each message*

*Register handler for
each message*

*Register other
channel events*

Create connections



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {  
  
    private void uponOutConnUp(OutConnectionUp ev, ...){  
        sendMessage(new PingMsg(currTime), ev.getPeer());  
    }  
  
    private void uponPing(PingMsg msg, Host from, ...){  
        sendMessage(new PongMsg(msg.getValue()), from);  
    }  
  
    private void uponPong(PongMsg msg, Host from, ...){  
        closeConnection(from);  
        openConnection(random(peers));  
    }  
  
}
```



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {
```

```
    private void uponOutConnUp(OutConnectionUp ev, ...){  
        sendMessage(new PingMsg(currTime), ev.getPeer());  
    }
```

```
    private void uponPing(PingMsg msg, Host from, ...){  
        sendMessage(new PongMsg(msg.getValue()), from);  
    }
```

```
    private void uponPong(PongMsg msg, Host from, ...){  
        closeConnection(from);  
        openConnection(random(peers));  
    }
```

```
}
```

*Handle channel
events*



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {
```

```
    private void uponOutConnUp(OutConnectionUp ev, ...){  
        sendMessage(new PingMsg(currTime), ev.getPeer());  
    }
```

```
    private void uponPing(PingMsg msg, Host from, ...){  
        sendMessage(new PongMsg(msg.getValue()), from);  
    }
```

```
    private void uponPong(PongMsg msg, Host from, ...){  
        closeConnection(from);  
        openConnection(random(peers));  
    }
```

```
}
```

*Handle channel
events*

Send messages



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {  
  
    private void uponOutConnUp(OutConnectionUp ev, ...){  
        sendMessage(new PingMsg(currTime), ev.getPeer());  
    }  
  
    private void uponPing(PingMsg msg, Host from, ...){  
        sendMessage(new PongMsg(msg.getValue()), from);  
    }  
  
    private void uponPong(PongMsg msg, Host from, ...){  
        closeConnection(from);  
        openConnection(random(peers));  
    }  
  
}
```

*Handle channel
events*

Send messages

*Handle received
messages*



Babel: API - Networking

```
public class PingPongProto extends GenericProtocol {  
  
    private void uponOutConnUp(OutConnectionUp ev, ...){  
        sendMessage(new PingMsg(currTime), ev.getPeer());  
    }  
  
    private void uponPing(PingMsg msg, Host from, ...){  
        sendMessage(new PongMsg(msg.getValue()), from);  
    }  
  
    private void uponPong(PongMsg msg, Host from, ...){  
        closeConnection(from);  
        openConnection(random(peers));  
    }  
  
}
```

*Handle channel
events*

Send messages

*Handle received
messages*

*Open/close
connections*



Babel: Summary

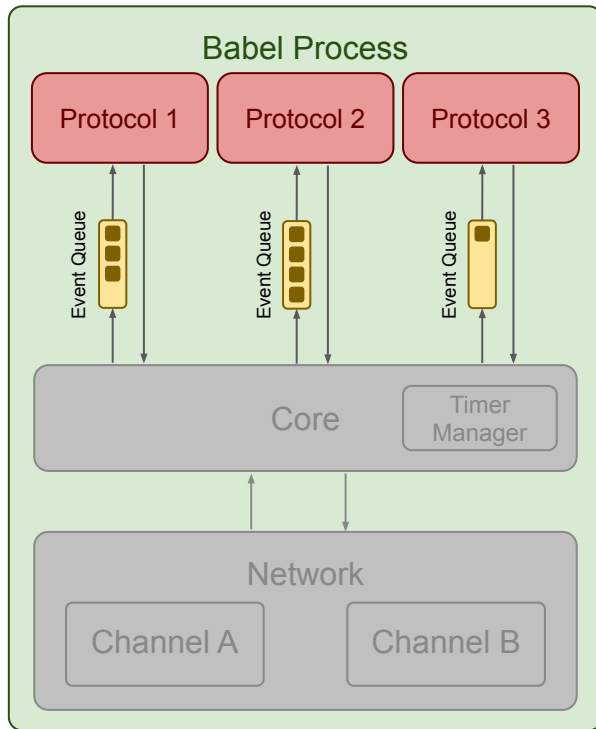
- **Developers** only implement **handlers for events**, and register them.
- Implementation **closely matches** typical algorithm specification (**pseudo-code**).
- **Shields developers** from the complexities of concurrency, multiplexing and networking
- Abstractions are sufficiently **generic** to support a wide variety of algorithms.
- Capable of supporting **production-ready performant** implementations.



Babel: Summary

- **Developers** only implement **handlers for events**, and register them.
- Implementation **closely matches** typical algorithm specification (**pseudo-code**).
- **Shields developers** from the complexities of concurrency, multiplexing and networking
- Abstractions are sufficiently **generic** to support a wide variety of algorithms.
- Capable of supporting **production-ready performant** implementations.

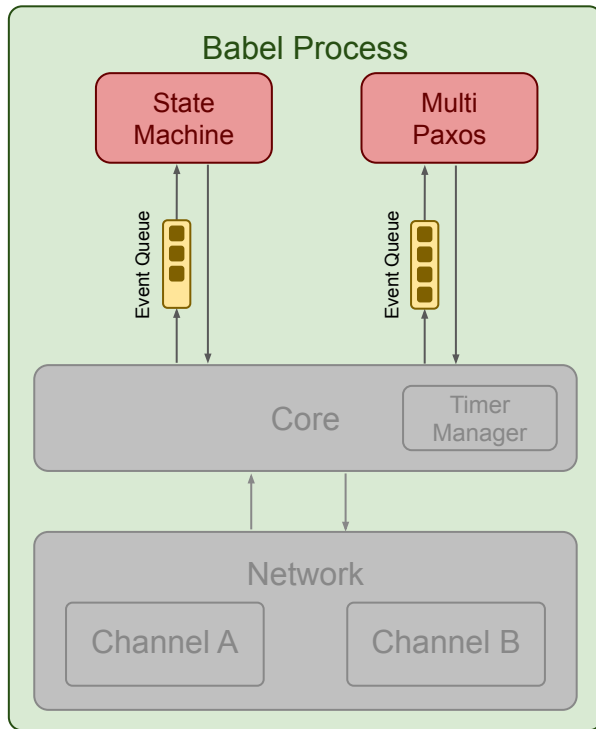
Class Example - MultiPaxos



Two Protocols

- State Machine
- MultiPaxos

Class Example - MultiPaxos

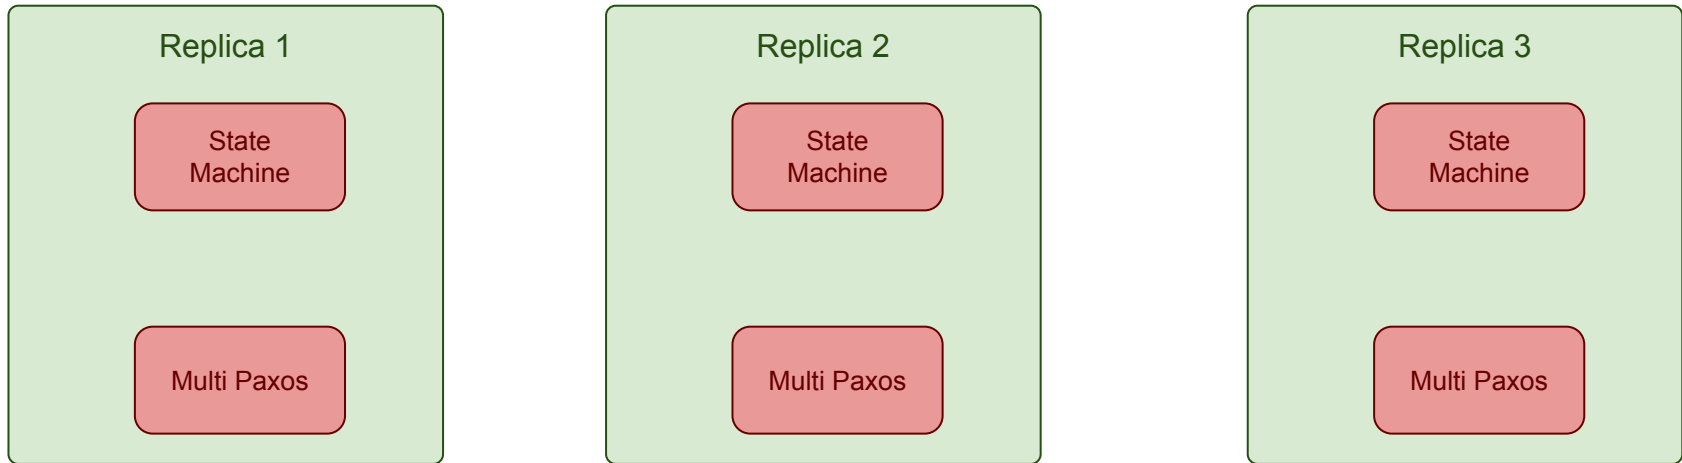


Two Protocols

- State Machine
- MultiPaxos

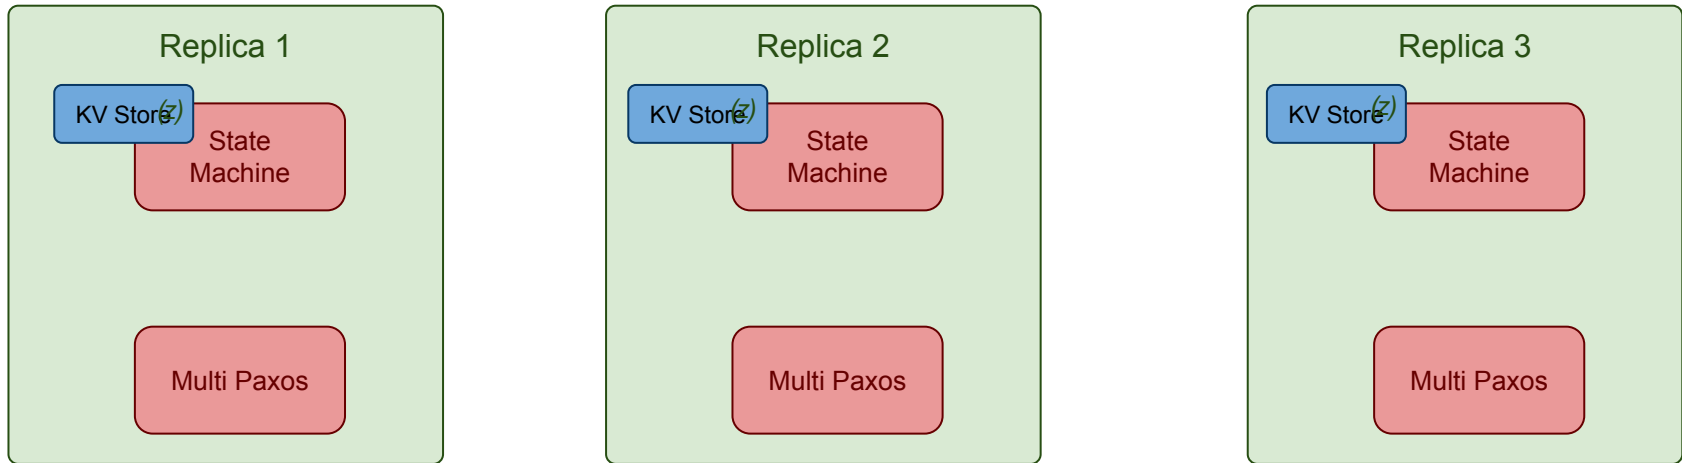


Class Example - MultiPaxos



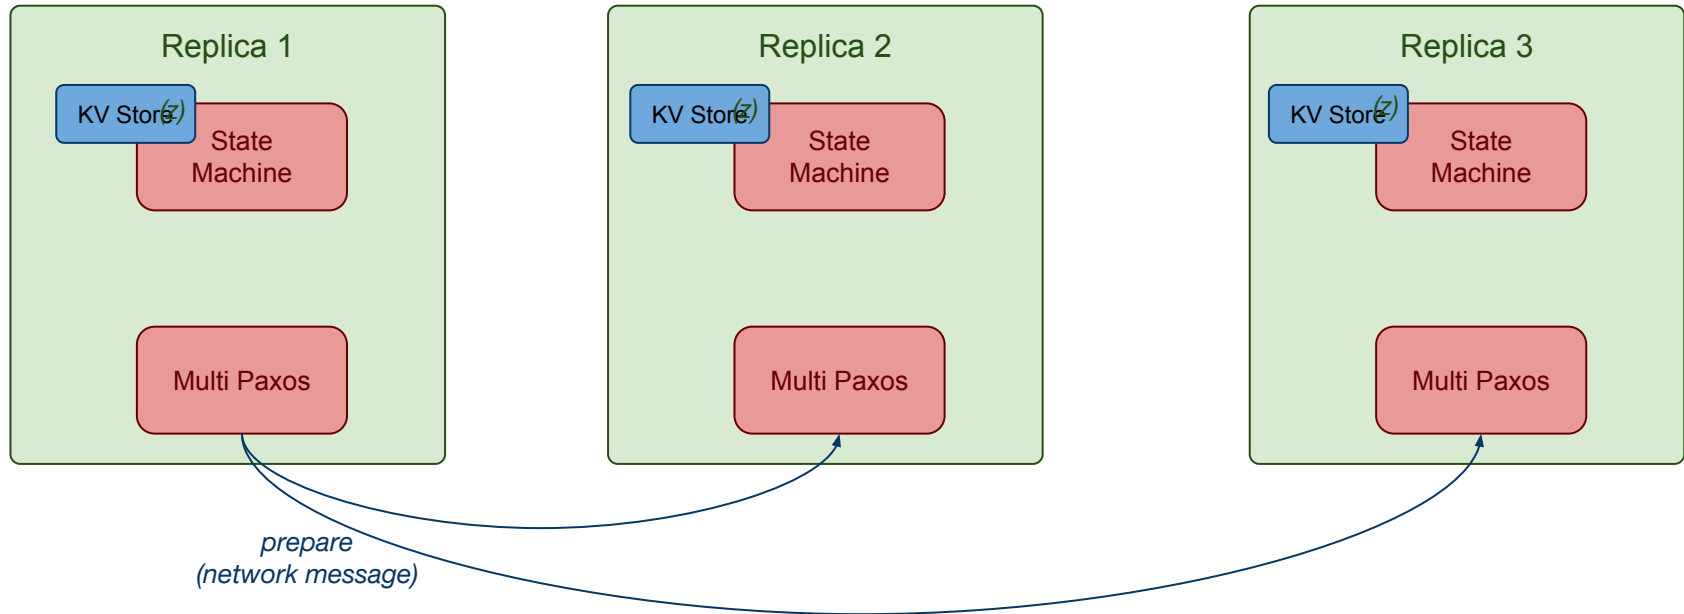


Class Example - MultiPaxos



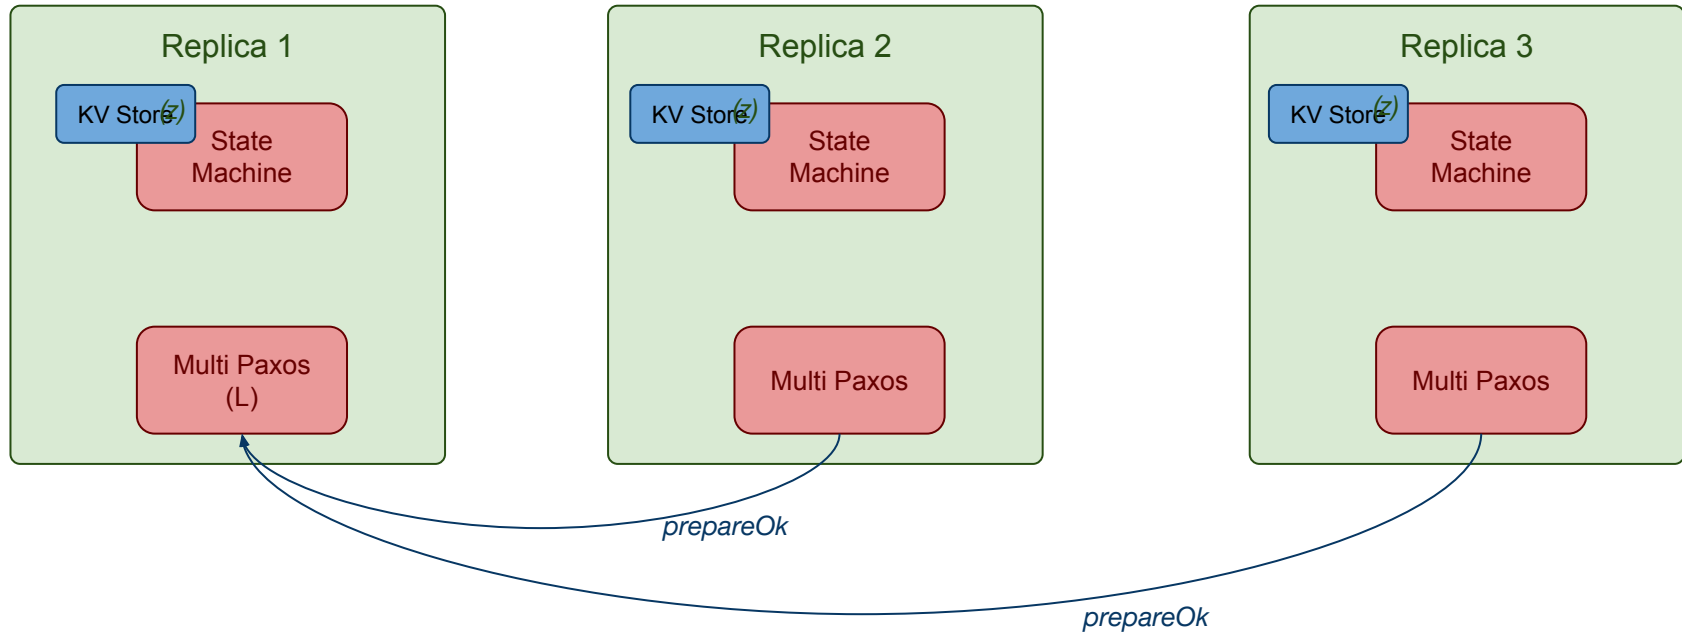


Class Example - MultiPaxos



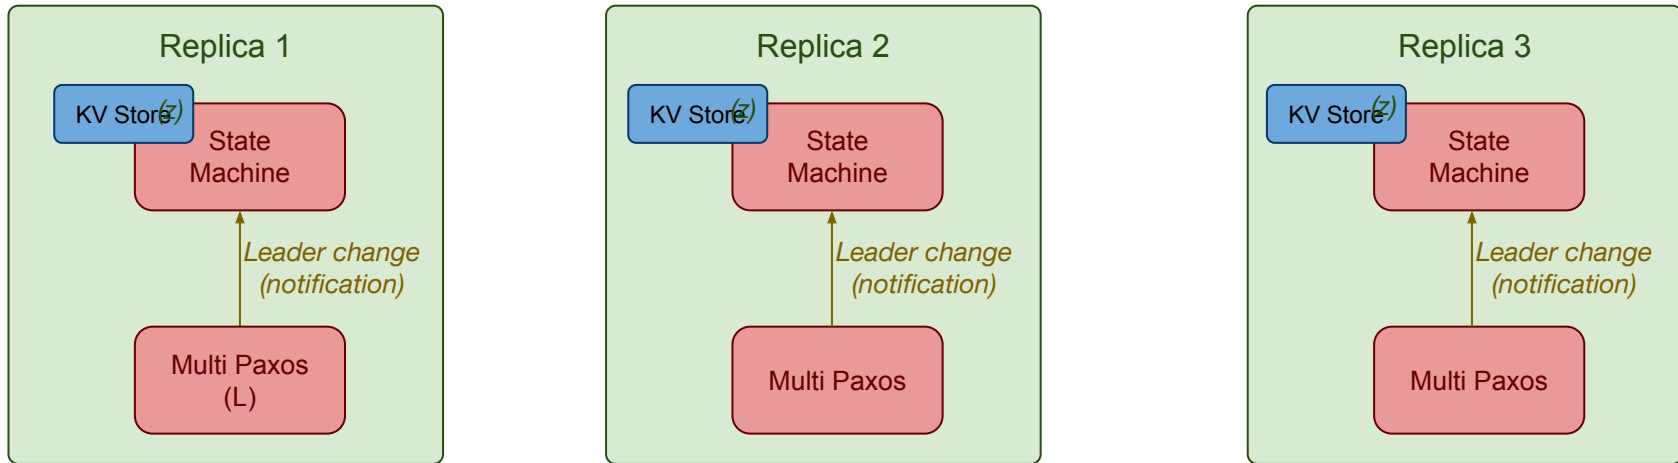


Class Example - MultiPaxos



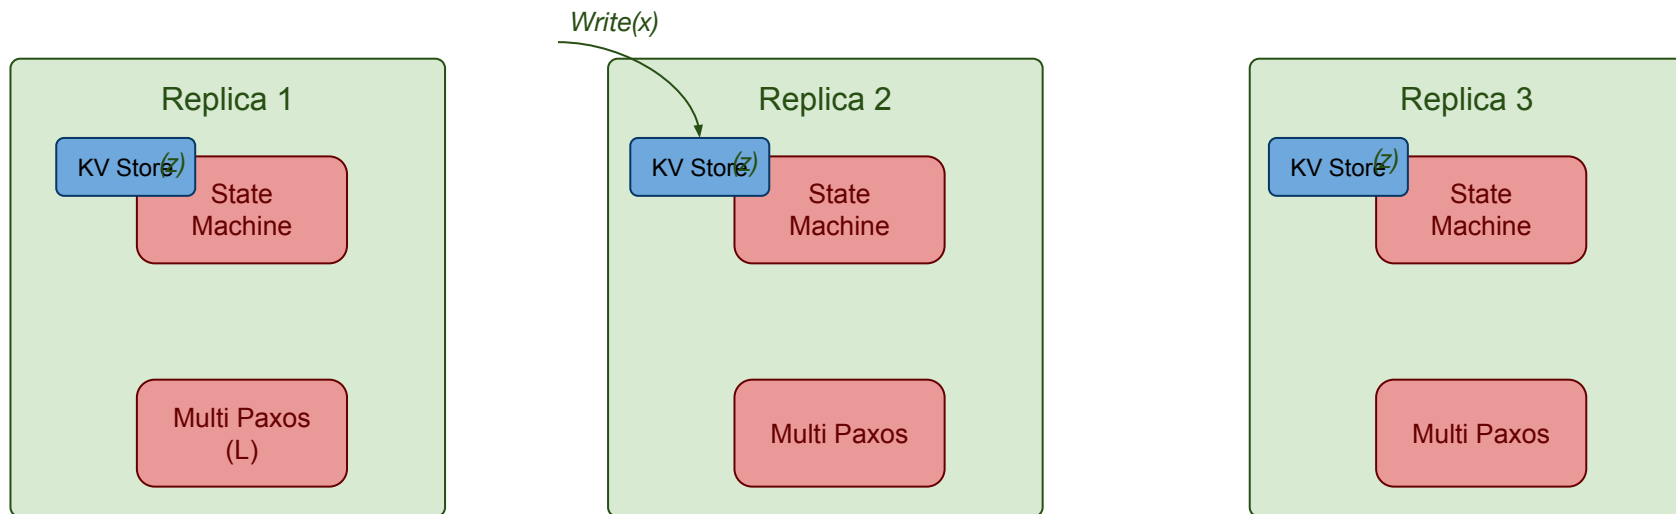


Class Example - MultiPaxos



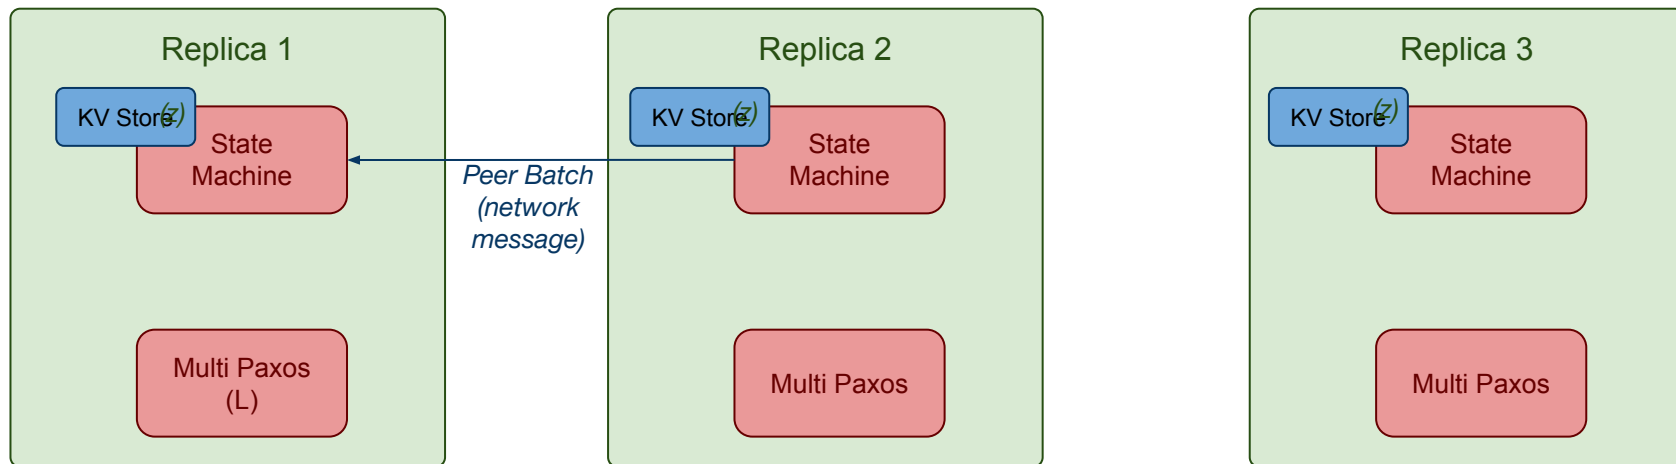


Class Example - MultiPaxos



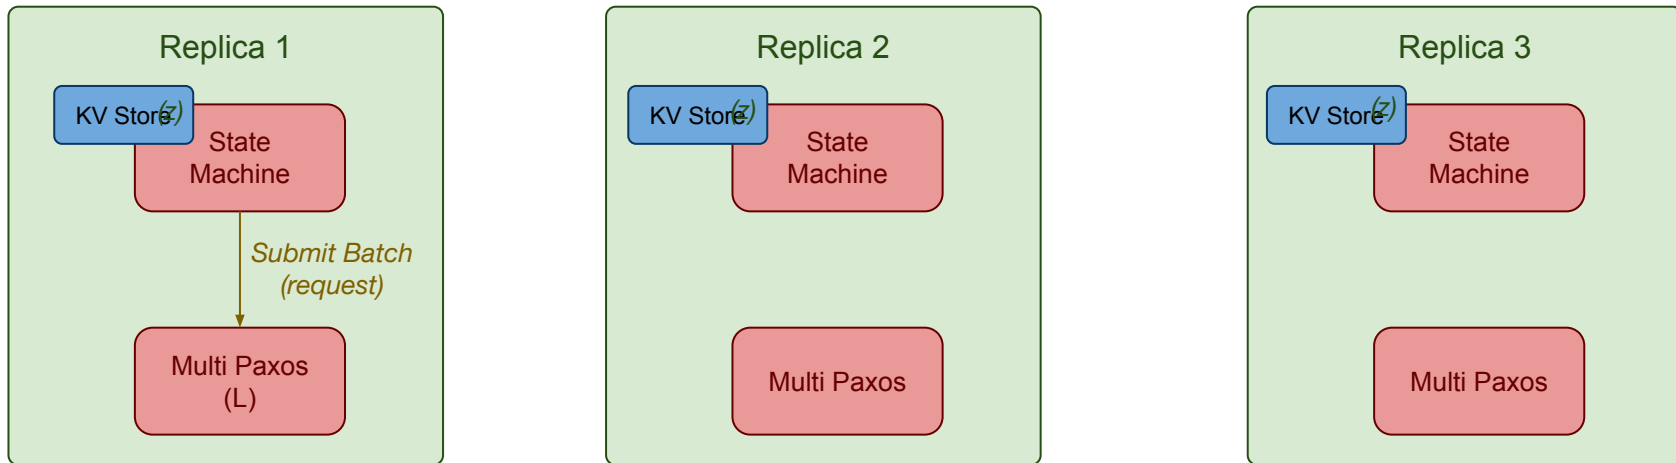


Class Example - MultiPaxos



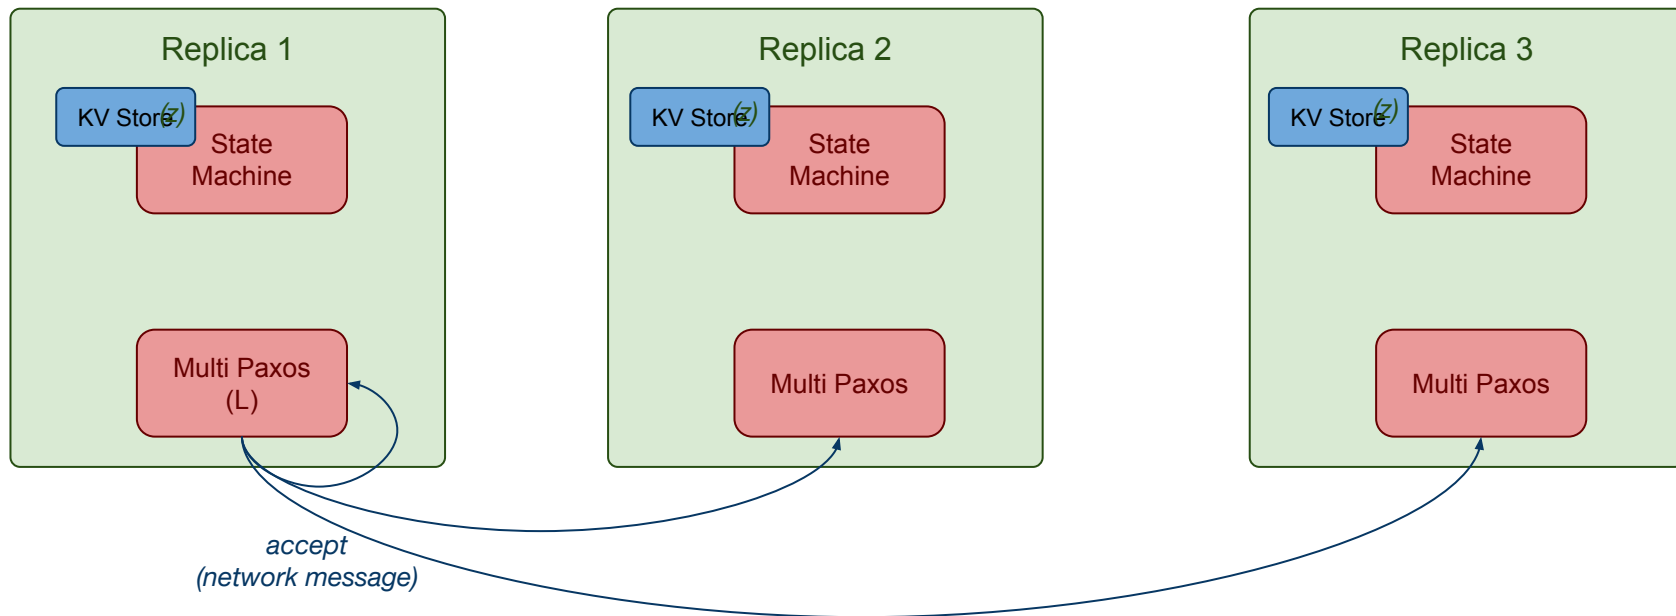


Class Example - MultiPaxos



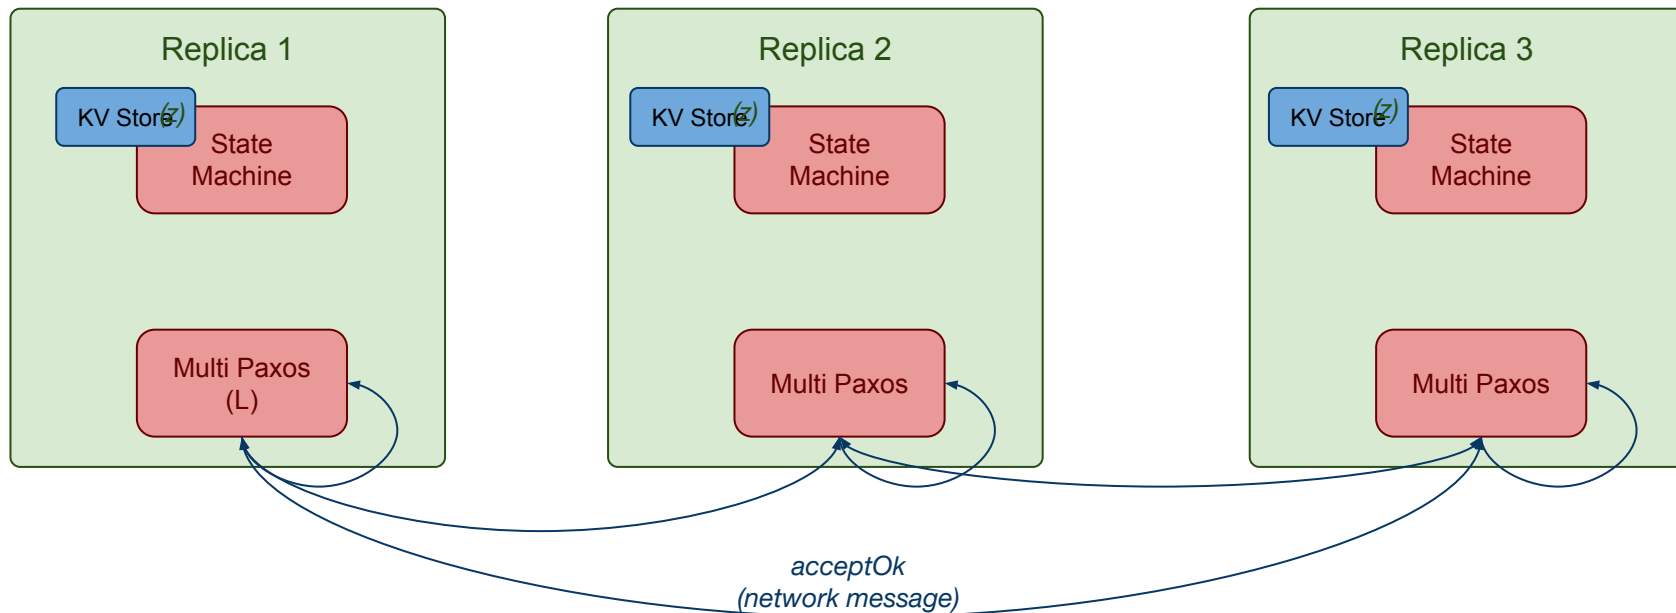


Class Example - MultiPaxos



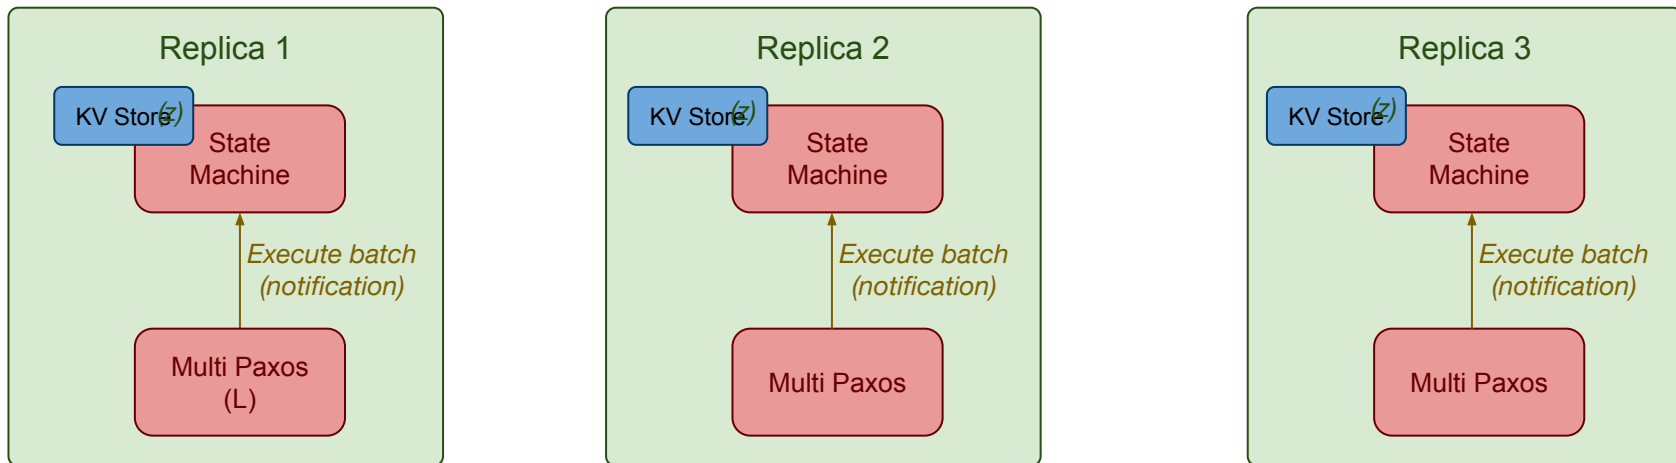


Class Example - MultiPaxos



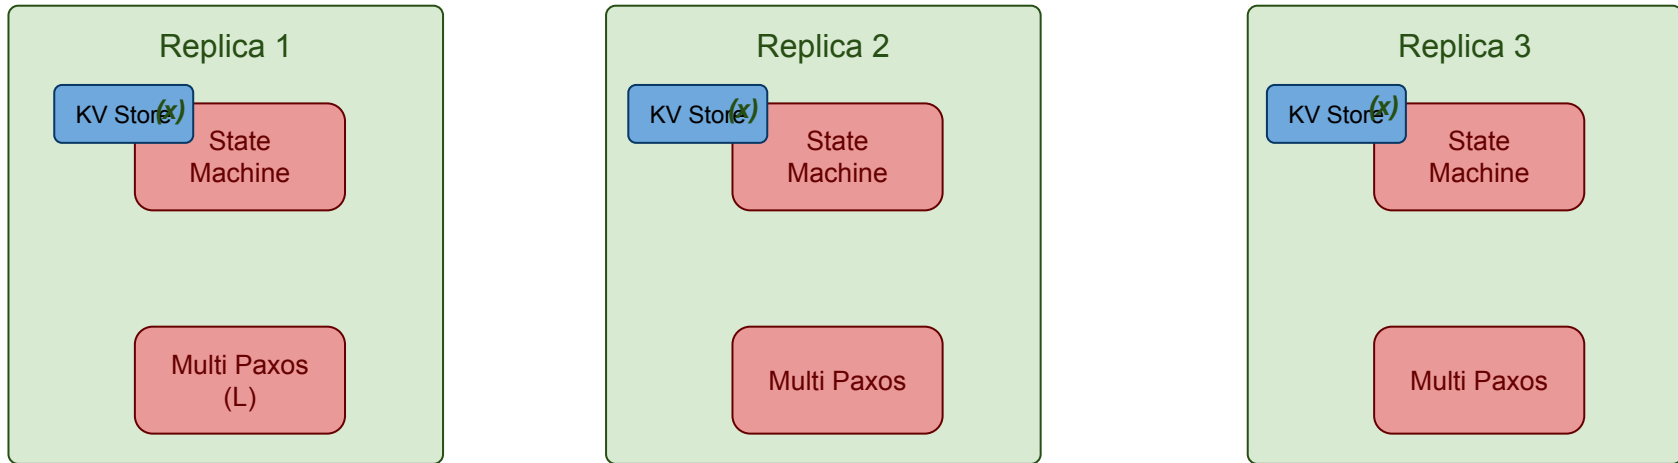


Class Example - MultiPaxos



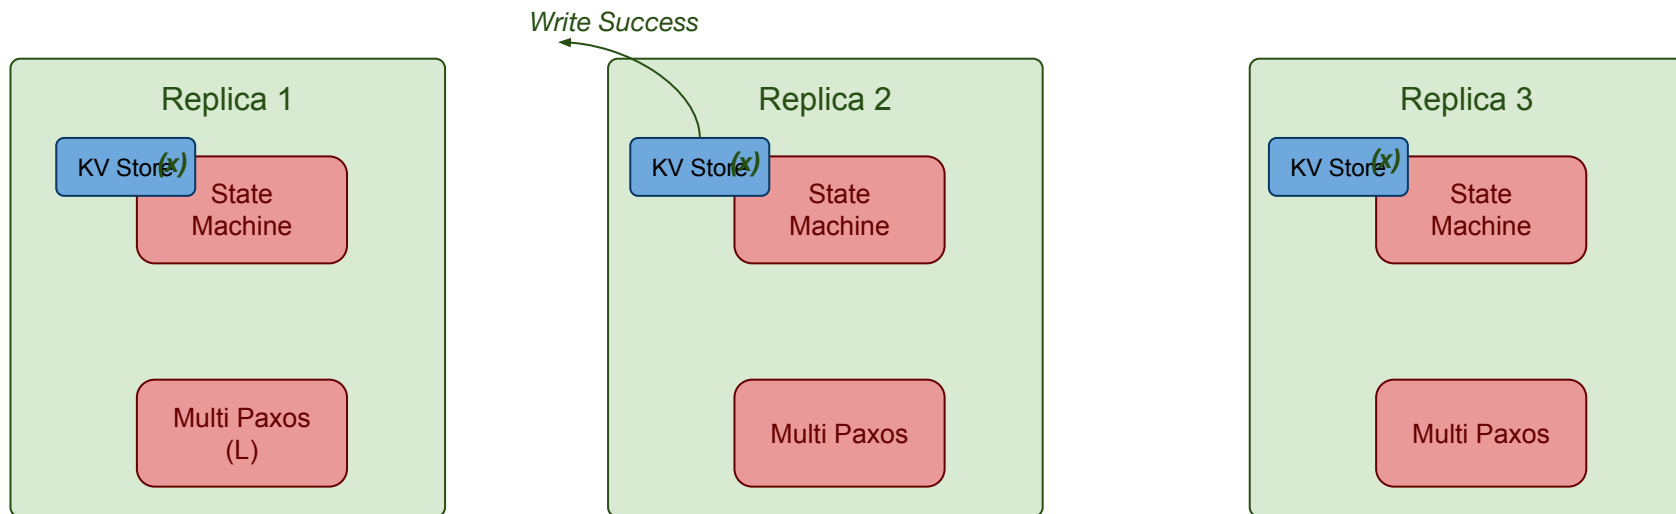


Class Example - MultiPaxos





Class Example - MultiPaxos





Babel: A framework for developing performant and dependable distributed protocols

Pedro Fouto, Pedro Ákos Costa,
Nuno Preguiça, João Leitão

SRDS 2022

<https://github.com/pfouto/babel-core>

<https://github.com/pedroAkos/babel-case-studies>

